

BUSINESS PLAN · CONFIDENTIAL

The intelligence layer for global drug supply.

Mederti turns the world's fragmented drug-shortage and recall data into a single, searchable, AI-native source of truth — for the pharmacists, hospitals, regulators and suppliers who keep medicine on the shelf.

29k+

Live shortage
events tracked

20+

National regulators
ingested daily

68

Automated data
scrapers

Live

Shipping real data
today

Prepared

June 2026

Stage

Live product · Pre-seed

Product

mederti.vercel.app

Category

Healthcare data & AI intelligence

A real, growing, expensive problem — with no system of record.

Drug shortages have roughly tripled over the past decade. Every developed health system now runs a permanent shortage desk, yet the underlying data lives in dozens of incompatible national registries, PDFs and email alerts. No one owns the global picture. Mederti does.

The problem

A hospital pharmacist learning that a critical injectable is short must today check the FDA, their national regulator, wholesaler portals and manufacturer notices — each in a different format, none cross-referenced, none predictive. The answer to "what do I do about it" doesn't exist in one place.

The solution

Mederti ingests every major regulator daily, resolves drugs to a canonical identity across brands, salts and countries, and layers Claude-level synthesis on top. One search returns the global shortage status, alternatives, recalls and an operational answer — grounded in the database, not hallucinated.

Why now

Three forces converge: shortages at record highs and politically salient; regulators mandating supply-chain transparency (EU, US, UK); and frontier LLMs finally able to synthesise messy regulatory text reliably. The data was always there — the means to unify and reason over it arrived in 2024–25.

Where we are

Not a deck — a **live product**. 29,000+ shortage events, 20+ regulators on daily cron, persona-aware UX for pharmacists, procurement and suppliers, and a tool-using Claude chat grounded in the database. Built and shipping.

\$3–4B

Est. annual cost of shortage management to US hospitals alone (labour + substitution)

\$2B+

Serviceable software & data spend across pharmacy, procurement & supplier intelligence

15M+

Potential professional end-users across target personas globally

The ask

We are raising a **pre-seed round** to convert a working product into a commercial platform: harden the data pipeline off laptop-cron onto cloud infrastructure, ship the institutional API and alerting tier, and land the first 10 paying hospital and supplier customers. Detail in §9.

The data exists. The intelligence doesn't.

Drug shortages are a structural feature of modern pharmaceutical supply chains — and the people managing them are flying blind across a dozen disconnected sources.

Fragmented

Every country runs its own shortage registry — FDA, TGA, EMA, MHRA, Health Canada, plus 15+ others — in different schemas, languages and update cadences. None talk to each other.

Unstructured

The same molecule appears as a brand, a salt form, an INN and a local trade name. "Is amoxicillin short?" has no single answer because the data was never resolved to one entity.

Backward-looking

Registries tell you a drug *is* short. None tell you what to do, what the alternatives are, or which drug is *about* to go short next.

Who feels the pain — and what it costs them

PERSONA	TODAY'S WORKFLOW	COST OF GETTING IT WRONG
Hospital pharmacist	Manually checks 4–6 portals per drug; rebuilds substitution plans by hand	Hours/day of clinical time; delayed or unsafe substitutions
Procurement / supply chain	Reacts to stockouts after they happen; no forward signal	Emergency spot-buys at 5–50× price; tendering chaos
Suppliers & manufacturers	No visibility into where demand gaps are opening globally	Missed revenue; mis-allocated production
Regulators & government	Own national data only; no cross-border early warning	Public-health exposure; political risk

"There is no Bloomberg Terminal for drug supply. Every stakeholder is reconstructing the same global picture, badly, every single day."

This is not a niche inconvenience. The U.S. alone consistently tracks **250–300+ active national shortages** at any time; the EU and individual member states track hundreds more. Managing them is a permanent, staffed, expensive function at every major health system on earth — and the tooling is spreadsheets and email.

One source of truth, with Claude on top.

Mederti unifies global regulatory data into a canonical drug graph, then makes it answerable in plain language — grounded in the database, never invented.

1 • Ingest everything, daily

68 automated scrapers pull shortage and recall data from 20+ national regulators every day. Each record is de-duplicated and normalised on the way in.

2 • Resolve to canonical identity

Brands, salt forms, synonyms and trade names roll up to a single INN via RxNorm and WHO ATC mappings — so *LORSTAT* and *atorvastatin* are one answer, not twelve.

3 • Layer live intelligence

Recalls link causally to shortages. Peer-set lead-time analysis across 16 EU markets surfaces forward signals. Pricing, alternatives and supplier data attach to every drug.

4 • Answer in plain language

A tool-using Claude agent queries the live database — 29k+ events and a 124-entry intelligence catalogue — and returns an operational answer with sources. Macro questions it can't ground, it declines rather than fabricates.

ONE QUERY • ONE ANSWER

"Is amoxicillin short, and what do I switch to?"

- ✓ 17 active shortages resolved across markets
- ✓ Canonical INN — every brand & salt rolled up
- ✓ Alternatives ranked, with their own status
- ✓ Linked recalls flagged as causal precursors
- ✓ A grounded, sourced operational answer

Live today at mederti.vercel.app — not a mockup.

Persona-aware by design

The same data renders differently for each user. A pharmacist sees a radically simplified *answer + actions* card. Procurement sees supply-risk and sourcing. Suppliers see demand gaps and market opportunities. One backend, four products.

A category waiting to be defined.

There is no incumbent "system of record" for global drug supply. The spend exists — scattered across point tools, manual labour and emergency procurement — but no one has unified it.

\$18B+

TAM — global pharma supply-chain & data intelligence software

\$2.4B

SAM — shortage/recall intelligence for pharmacy, procurement & suppliers

\$120M

SOM — beachhead: hospital pharmacy + supplier intel, English-speaking + EU markets, yrs 1–3

TAM/SAM/SOM are directional estimates for planning, triangulated from health-system count, pharmacy staffing and comparable data-platform pricing — not audited market figures.

Multiple, stacking buyers

- **Hospitals & health systems** — ~7,000 hospitals in the US/EU/UK/ANZ alone, each with a shortage function
- **Pharmacy networks** — chains and independents needing live substitution data
- **Manufacturers & suppliers** — global demand-gap intelligence to allocate production
- **Regulators & government** — cross-border early-warning no single agency can build alone
- **Researchers & press** — citation-grade data = top-of-funnel moat

Why this market opens now

- **Shortages at record highs** — politically and clinically urgent worldwide
- **Regulatory mandates** — EU/US/UK forcing supply-chain transparency and reporting
- **LLMs crossed a threshold** — synthesising messy regulatory text is now reliable
- **Geopolitical supply risk** — API concentration in a few countries is now a board-level concern

The wedge: own the data and the canonical drug graph first — the asset no competitor and no single regulator has. Everything else (alerting, API, supplier marketplace, forecasting) is a product layer on a moat that compounds daily.

Already built. Already live.

Most plans at this stage are promises. Mederti is shipping real data to a live URL today — the platform, pipeline and AI layer already exist.

What's live SHIPPING

- ✓ 29,000+ shortage events across 20+ regulators
- ✓ 68 scrapers, idempotent daily ingestion
- ✓ Canonical drug graph — 16,000+ brand→INN links
- ✓ Persona-aware drug pages (pharmacist / procurement / supplier)
- ✓ Tool-using Claude chat, grounded in the database
- ✓ Recall→shortage causal linking
- ✓ Peer-set predictive signals across 16 EU markets
- ✓ Email alerts & watchlists

Near-term roadmap BUILDING

- Migrate pipeline off laptop-cron to cloud (reliability)
- Institutional **API tier** & bulk data exports
- Forward-signal model — which drug shortages *next*
- Wire 12+ built-but-unscheduled country scrapers (India, China, Brazil...)
- Supplier marketplace & demand-gap matching
- Public methodology & data-freshness dashboard

The architecture is real and modern

LAYER	STACK	STATE
Data ingestion	Python 3.11, 68 scrapers, MD5 de-dup	LIVE
Database	Supabase / PostgreSQL — 46 migrations, RLS, full-text + trigram search	LIVE
Application	Next.js 16, React 19, 30+ API route handlers	LIVE
AI layer	Claude (tool-using) with DB grounding + web search	LIVE
Infra hardening	Cloud cron migration, monitoring, freshness SLAs	IN PROGRESS

Translation: the capital-intensive, slow part — building the data asset and the platform — is **done and running**. Investment accelerates commercialisation, not basic construction.

Data underneath, four ways to monetise on top.

A single compounding data asset supports recurring revenue across every persona — classic data-platform economics with high gross margin and strong retention.

TIER	BUYER	WHAT THEY GET	PRICING MODEL	ACV (ILLUSTRATIVE)
Free / Pro	Individual pharmacists, clinicians	Search, drug pages, alerts, AI chat	Freemium → seat	\$0–300
Institutional	Hospitals, pharmacy networks	Team seats, watchlists, substitution workflows, SSO	Per-site SaaS	\$8–40k
Supplier Intelligence	Manufacturers, distributors	Demand-gap signals, market opportunities, marketplace	SaaS + leads	\$25–100k
Data & API	Govt, regulators, researchers, enterprise	Bulk exports, API access, custom feeds	Data licence	\$30–250k

High gross margin

Data is collected once, sold many times. Marginal cost per customer is near-zero — the pipeline runs regardless of seat count.

Compounding moat

Every day of ingestion deepens the historical record and the canonical graph — an asset a new entrant cannot backfill.

Sticky by workflow

Once a shortage desk runs on Mederti, switching means rebuilding a daily-critical workflow. Retention is structural.

Land & expand: free pharmacist search seeds organic, bottom-up adoption inside hospitals → institutional contracts → supplier and data-licence revenue layered on the same backend. The cheapest user to acquire funds the most valuable one to sell.

Fragments of the answer exist. The whole doesn't.

Today's "competitors" each own one slice — one country, one persona, or static data with no intelligence. None unify global data with AI synthesis.

CATEGORY	EXAMPLES	THEIR LIMIT
National registries	FDA, EMA, TGA databases	Single-country, raw, unresolved, no synthesis or alternatives
Pharmacy point tools	ASHP, wholesaler portals	US-centric, manual, no cross-border or supplier view
Supply-chain suites	Enterprise procurement ERPs	Internal-inventory focused, blind to global regulatory signal
Generic AI chat	ChatGPT, web search	No grounded drug-supply data; hallucinates status & dates

Why Mederti wins — and keeps winning

- **The canonical drug graph**

Resolving brands, salts, synonyms and trade names to one INN across countries is the hard, unglamorous asset competitors haven't built. It's what makes one query return one answer.
- **Daily compounding history**

29k+ events and growing every day. The longitudinal record — which drug shorted, when, why, for how long — cannot be retro-acquired. Time is the moat.
- **AI-native, grounded**

Claude synthesis sits on real data with guardrails against fabrication — the trust bar healthcare demands. Generic chat can't, and registries won't.
- **Multi-persona network**

Pharmacists' demand signals become suppliers' market intelligence. Each side makes the other more valuable — a two-sided data network on one backend.

Bottom-up adoption, top-down contracts.

Win the individual pharmacist for free; convert their institution; expand into suppliers and data licensing. Credibility content seeds the top of the funnel.

Phase 1 — Seed

Land the practitioner. Free search + AI chat for pharmacists and clinicians. SEO & citation-grade data make Mederti the default answer to "is X short?"

Phase 2 — Convert

Sell the institution. Bottom-up usage inside hospitals becomes the wedge for per-site contracts with team workflows, SSO and SLAs.

Phase 3 — Expand

Monetise the network. Supplier intelligence and government/enterprise data licences on the same compounding backend.

Growth engines

- **Content & SEO moat** — a public, cited methodology and freshness dashboard make Mederti the reference researchers and press link to
- **Product-led free tier** — zero-friction entry for the highest-volume persona
- **Regulatory tailwind** — transparency mandates create budget and urgency at institutions
- **Design partnerships** — a handful of lighthouse hospitals & suppliers shape the institutional product
- **Supplier pull** — manufacturers pay for demand-gap signal they can't get elsewhere
- **Geographic expansion** — each new country scraper is a new addressable market at near-zero marginal cost

First commercial milestone: 10 paying institutional/supplier logos within 12 months of funding — enough to validate ACV, prove retention, and anchor the seed round.

Lean to build, leveraged to scale.

Because the platform already exists, capital funds commercialisation — not construction. Data-platform margins mean revenue compounds far faster than cost.

Illustrative 3-year trajectory

METRIC	YEAR 1	YEAR 2	YEAR 3
Paying institutional / supplier logos	10	45	140
Blended ACV	\$22k	\$30k	\$38k
ARR	\$0.2M	\$1.4M	\$5.3M
Gross margin	~80%	~85%	~88%
Free-tier users (top of funnel)	25k	120k	400k

Illustrative model for planning purposes, not a forecast or guarantee. Assumes pre-seed close in 2026, cloud-pipeline migration complete by Q1, and API/institutional tiers live by mid-year.

Cost structure

- **Engineering & data** — pipeline reliability, API, forecasting model
- **Cloud & AI inference** — scales with usage, not headcount
- **Go-to-market** — founder-led sales + content, then first AE

Unit economics drivers

- Near-zero marginal cost per customer (data collected once)
- Workflow lock-in → low churn, high net revenue retention
- Free tier as a low-CAC engine for institutional conversion

Fund the path from product to platform.

We are raising a pre-seed round to harden infrastructure, ship the commercial tiers, and land the first paying customers — turning a live product into a revenue-generating platform.

40% · Engineering & Data

Cloud pipeline migration, institutional API, forward-signal model, +12 country scrapers.

35% · Go-to-Market

Founder-led sales, design partnerships, content/SEO engine, first commercial hire.

25% · Operations

Compliance & data governance, monitoring & SLAs, runway to commercial milestones.

Milestones this round buys

Quarter 1

Reliability. Pipeline off laptop-cron onto cloud infra; monitoring, freshness SLAs and a public data-freshness dashboard live.

Quarter 2

Commercial tiers. Institutional seats + API/data-export tier shipped; first 3 design-partner logos signed.

Quarter 3

Coverage & signal. 12+ new country scrapers live; forward-signal model in beta with lighthouse customers.

Quarter 4

Validation. 10 paying logos, proven ACV & retention, supplier-intelligence revenue line open — seed-ready.

WHY BACK MEDERTI NOW

Built

Live product & real data — capital accelerates, not constructs.

Timely

Record shortages + regulatory mandates + LLMs that finally work.

Defensible

A canonical data asset that compounds daily and can't be backfilled.



The world's leading source of drug-supply intelligence.

The data was always there, scattered and unusable. Mederti makes it one source of truth — and puts Claude-level reasoning on top. The product is live. The opportunity is open. The moat compounds every day we run.

Live product

mederti.vercel.app

Stage

Pre-seed

Contact

robfindlay@gmail.com